




Practical Concurrent and Parallel Programming III

Performance Measurements

Jørgen Staunstrup

Motivations for Concurrency




Inherent: User interfaces and other kinds of input/output

Exploitation: Hardware capable of simultaneously executing multiple streams of statements

Hidden: Enabling several programs to share some resources in a manner where each can act as if they had sole ownership

Motivation for performance measurements - 1



Why is creating a Thread said to be expensive?

▲ 0
▼ votes
1 view

Asked Oct 29, 2019 in [Java](#) by [Anvi](#) (10.2k points)

The Java tutorials say that creating a Thread is expensive. But why exactly is it expensive? What exactly is happening when a Java Thread is created that makes its creation expensive? I'm taking the statement as true, but I'm just interested in mechanics of Thread creation in JVM.

Threads are expensive

But how expensive ?


~600 ns to create (on this laptop)

~20 times more time than creating a simple object

40000 ns to start a thread !!! (on this laptop)

Today: How to get such numbers !

Motivation for performance measurements - 2



Sorting 1_000_000 numbers with Quicksort takes 9 μ s (using 1 Thread)

How fast can we do it using N threads?

www.menti.com : 38 23 09 6

Explanation and animation:

<https://en.wikipedia.org/wiki/Quicksort>

Performance measurements: motivation and introduction

Pitfalls (and avoiding them)

Calculating means and variance (efficiently)

Measurements of thread and lock overhead

Algorithms for parallel computing

Performance measurements: motivation and introduction

Pitfalls (and avoiding them)

Calculating means and variance (efficiently)

Measurements of thread and lock overhead

Algorithms for parallel computing

(Performance) Measurements



Key in many sciences (experiments, observations, predictions, ...)

A bit of statistics

A bit of numerical analysis

A bit of computer architecture (number representation)

Code for measuring execution time

Based on Microbenchmarks in Java and C# by Peter Sestoft (see [benchmarkingNotes.pdf](#) in material for this week)

All numbers in these slides were measured in August 2021 on a:

Intel Core i5-1035G4 CPU @ 1.10GHz, 4 Core(s), 8 Logical Processor(s)

Example: measuring a (simple) function



```
private static int multiply(int i) {  
    return i * i;  
}
```

```
start= System.nanoTime();  
multiply(126465);  
end= System.nanoTime();  
  
System.out.println(end-start+" ns");
```

1700 ns

1500 ns

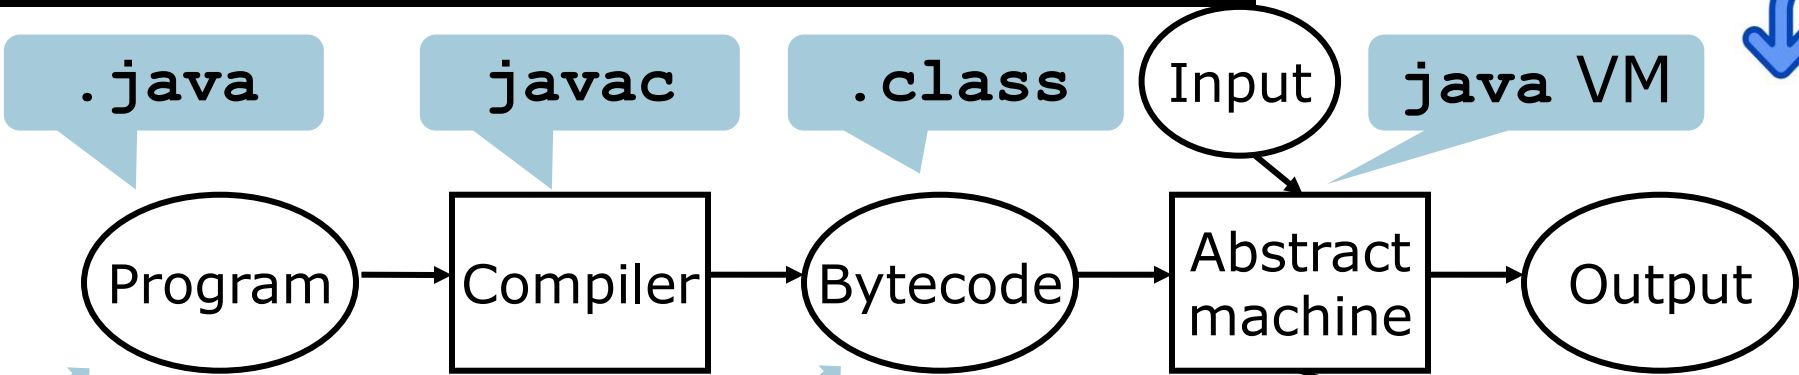

2500 ns

$\sim 1 - 2\text{ns}$

What is going on?

Java compiler and virtual machine

9



```
for (int i=0; i<n; i++)  
    sum += sqrt(arr[i]);
```

```
21 iconst_0  
22 istore 5  
24 iload 5  
26 iload 2  
27 if_icmpge      46  
30 dload 3  
31 aload 1  
32 iload 5  
34 daload  
35 invokestatic  Math.sqrt: (D)D  
38 dadd  
39 dstore 3  
40 iinc 5, 1  
43 goto 24
```

JVM

JIT (Just In Time)

```
19 xorl %ebx,%ebx  
1b jmp 3a  
1d leal 0x00(%ebp),%ebp  
20 fldl 0xec(%ebp)  
23 cmpl %ebx,0xc(%edi)  
26 jbe 49  
2c leal  
0x10(%edi,%ebx,8),%eax  
...
```

x86



Performance measurements: motivation and introduction

Pitfalls (and avoiding them)

Calculating means and variance (efficiently)

Measurements of thread and lock overhead

Algorithms for parallel computing



Microbenchmarks in Java and C#

Peter Sestoft (sestoft@itu.dk)

IT University of Copenhagen, Denmark

Version 0.8.0 of 2015-09-16

A goldmine of good advice

Accompanying code: [Benchmark.java](#)

Abstract: Sometimes one wants to measure the speed of software, for instance, to measure whether a

```
class Benchmark {
    public static void main(String[] args) { new Benchmark(); }

    public Benchmark() {
        SystemInfo();
        // Mark0();
        // Mark1();
        Mark2();
        ...
        // Mark8("random_index", i -> rnd.nextInt(n));
        ...
        // SortingBenchmarks();
        ...
    }
}
```

Example: measuring a simple function

12



```
private static double multiply(int i) {  
    double x = 1.1 * (double)(i & 0xFF);  
    return x * x * x * x * x * x * x * x * x * x * x * x *  
           * x * x * x * x * x * x * x * x * x * x * x * x;  
}
```

```
public static double Mark2() {  
    Timer t = new Timer();  
    int count = 100_000_000;  
    double dummy = 0.0;  
    for (int i=0; i<count; i++)  
        dummy += multiply(i);  
    double time = t.check() * 1e9 / count;  
    System.out.printf("%6.1f ns%n", time);  
    return dummy;  
}
```

Get the code from
timingMultiplication.java
Try running it yourself
Report result in poll
www.menti.com : 38 23 09 6

```
# OS:    Windows 10; 10.0; amd64  
# JVM:   Oracle Corporation; 1.8.0_181  
# CPU:   Intel64 Family 6 Model 126 Stepping 5, GenuineIntel; 8 "cores"  
# Date:  2021-09-12T09:14:34+0200  
24.0 ns
```

The Timer class (in Benchmark.java)

13



A simple Timer class for Java

Works on all platforms (Linux, MacOS, Windows)

```
public class Timer {  
    private long start, spent = 0;  
    public Timer() { play(); }  
    public double check()  
    { return (System.nanoTime()-start+spent)/1e9; }  
    public void pause() { spent += System.nanoTime()-start; }  
    public void play() { start = System.nanoTime(); }  
}
```

Automating multiple runs (Mark3)



Results will usually vary

```
public static double Mark3() {  
    int n = 10;  
    int count = 100_000_000;  
    double dummy = 0.0;  
    for (int j=0; j<n; j++) {  
        Timer t = new Timer();  
        for (int i=0; i<count; i++)  
            dummy += multiply(i);  
        double time = t.check() * 1e9 / count;  
        System.out.printf("%6.1f ns%n", time);  
    }  
    return dummy;  
}
```

```
24.6 ns  
24.6 ns  
24.5 ns  
24.6 ns  
24.4 ns  
24.3 ns  
24.5 ns  
24.4 ns  
24.7 ns  
24.6 ns
```

What is the running time?



What should you report as the result, when the observations are:

30.7 ns 30.3 ns 30.1 ns 30.7 ns 30.5 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ??

Mean: 30.4 ns

What if they are:

30.7 ns 100.2 ns 30.1 ns 30.7 ns 20.2 ns 30.4 ns 2.0 ns 30.3 ns 30.5 ns 5.4 ns ??

Mean: 31.0 ns

Standard deviation/variance

16



$$\mu = \frac{1}{n} \sum_{j=1}^n t_j$$

Benchmark note p6

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j - \mu)^2}$$


30.7 ns 30.3 ns 30.1 ns 30.7 ns 50.2 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ??

Mean: 32.5 ns Standard deviation: 6.2

Normal distribution



17





What should you report as the result, when the observations are:

30.7 ns 30.3 ns 30.1 ns 30.7 ns 50.2 ns 30.4 ns 30.9 ns 30.3 ns 30.5 ns 30.8 ns ??

Mean: 32.5 ns Standard deviation: 6.2

50.2 is an outlier

because there is a probability of less than 4.6 % that 50.2 is a correct observation



Performance measurements: motivation and introduction

Pitfalls (and avoiding them)

Calculating means and variance (efficiently)

Measurements of thread and lock overhead

Algorithms for parallel computing

Computing the variance

20



$$\mu = \frac{1}{n} \sum_{j=1}^n t_j$$

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j - \mu)^2}$$

Requires two passes through the data

$$\sigma^2 = \frac{1}{n(n-1)} (n \sum_{j=1}^n t_j^2 - (\sum_{j=1}^n t_j)^2)$$

Can be done in one pass (on-line alg.)

```
for (int j=0; j<n; j++) {  
    Timer t = new Timer();  
    for (int i=0; i<count; i++)  
        ...  
    double time = t.check() * 1e9 / count;  
    st += time;  
    sst += time * time;  
}  
double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n) / (n-1));  
System.out.printf("%6.1f ns +/- %6.3f%n", mean, sdev);
```

The two formulas give the same result

21



$$\mu = \frac{1}{n} \sum_{j=1}^n t_j$$

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j - \mu)^2}$$

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{j=1}^n (t_j^2 + \mu^2 - 2t_j\mu)}$$

$$\sigma^2 = \frac{1}{n-1} \sum_{j=1}^n (t_j^2 + \mu^2 - 2t_j\mu)$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 + \sum_{j=1}^n (\mu^2 - 2t_j\mu))$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 + n\mu^2 - 2\mu \sum_{j=1}^n t_j)$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 + n\mu^2 - 2\mu n\mu)$$

$$\sigma^2 = \frac{1}{n-1} (\sum_{j=1}^n t_j^2 - n\mu^2)$$

$$\sigma^2 = \frac{1}{n(n-1)} (n \sum_{j=1}^n t_j^2 - \mu^2)$$

$$\sigma^2 = \frac{1}{n(n-1)} (n \sum_{j=1}^n t_j^2 - (\frac{1}{n} \sum_{j=1}^n t_j)^2)$$



Formula in Benchmark note

See exercises03.pdf

also https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance



Formula used in code (one pass algorithm)



$$\sigma^2 = \frac{1}{n(n-1)} \left(n \sum_{i=1}^n x_i^2 - \left(\sum_{i=1}^n x_i \right)^2 \right)$$

```
int n = 10;
...
for (int j=0; j<n; j++) {
    Timer t = new Timer();
    for (int i=0; i<count; i++)
        ...
    double time = t.check() * 1e9 / count;
    st += time;
    sst += time * time;
}
double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n) / (n-1));
System.out.printf("%6.1f ns +/- %6.3f%n", mean, sdev);
```

Beware: $sst - \text{mean} * \text{mean} * n$

can be a very small number

Beware of cancellation when subtracting numbers that are close to each other:

```
1010101000010110110001110101.111
-1010101000010110110001110001.100
-----
0000000000000000000000000100.011
```



Beware of cancellation when subtracting numbers that are close to each other:

```

1010101000010110110001110101.111
-1010101000010110110001110001.100
-----
0000000000000000000000000000100.011

```

(sst - mean*mean) can be problematic.

How to do it: https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance

Mark5 - computes mean and variance

25



```
public static double Mark5() {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j<n; j++) {
            Timer t = new Timer();
            for (int i=0; i<count; i++) dummy += multiply(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time;
            sst += time * time;
            totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
        System.out.printf("%6.1f ns +/- %8.2f %10d%n", mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}
```

Mark5 - computes mean and variance

26



```
public static double Mark5() {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j<n; j++) {
            Timer t = new Timer();
            for (int i=0; i<count; i++) dummy += multiply(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time;
            sst += time * time;
            totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
        System.out.printf("%6.1f ns +/- %8.2f %10d%n", mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}
```



```
private static double multiply(int i) {  
    . . .  
}
```

Java: `multiply(i)` is a number

Java: `i -> multiply(i)` is a function

<https://docs.oracle.com/javase/tutorial/java/javaOO/lambdaexpressions.html>

```
Mark6( . . . , i -> multiply(i));
```

www.menti.com : 38 23 09 6

Mark6 - introducing a functional argument

28



```
public static double Mark6(String msg, IntToDoubleFunction f) {
    int n = 10, count = 1, totalCount = 0;
    double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
    do {
        count *= 2;
        st = sst = 0.0;
        for (int j=0; j<n; j++) {
            Timer t = new Timer();
            for (int i=0; i<count; i++) dummy += f.applyAsDouble(i);
            runningTime = t.check();
            double time = runningTime * 1e9 / count;
            st += time; sst += time * time; totalCount += count;
        }
        double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
        System.out.printf("%-25s %15.1f ns %10.2f %10d%n", msg, mean, sdev, count);
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
    return dummy / totalCount;
}

public interface IntToDoubleFunction { double applyAsDouble(int i); }
```

The function **f** is benchmarked

```
Mark6("multiply", i -> multiply(i));
// same as line above, for motivation see here
// https://docs.oracle.com/javase/tutorial/java/javaOO/methodreferences.html
```


Example use of Mark6



```
Mark6("multiply", i -> multiply(i));
```

multiply	800.0	ns	1435.27	2
multiply	250.0	ns	0.00	4
multiply	212.5	ns	80.04	8
multiply	187.5	ns	39.53	16
multiply	200.0	ns	82.92	32
multiply	57.8	ns	24.26	64
multiply	46.9	ns	4.94	128
...				
multiply	30.6	ns	0.61	2097152
multiply	30.0	ns	0.10	4194304
multiply	30.1	ns	0.15	8388608

Mark7 - printing only final values



```
public static double Mark7(String msg, IntToDoubleFunction f) {  
    ...  
    do {  
        ...  
    } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);  
    double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));  
    System.out.printf("%-25s %15.1f %10.2f %10d%n", msg, mean, sdev, count);  
    return dummy / totalCount;  
}
```



Performance measurements: motivation and introduction

Pitfalls (and voiding them)

Calculating means and variance (efficiently)

Measurements of thread and lock overhead

Algorithms for parallel computing

Thread creation

32



```
Mark7("Thread create",
  i -> {
    Thread t = new Thread(() -> {
      for (int j=0; j<1000; j++) // not executed
        ai.getAndIncrement(); // thread t created, but not started
    });
    return t.hashCode();
  });
```

Takes 700 ns

Slow or fast?

Creating an object

33



A thread is an object, so let us start finding the cost of creating a simple object.

```
class Point {  
    public final int x, y;  
    public Point(int x, int y) { this.x = x; this.y = y; }  
}
```

```
Mark7("hashCode()", i -> myPoint.hashCode());
```

```
Mark7("Point creation",  
    i -> {  
        Point p = new Point(i, i);  
        return p.hashCode();  
    });
```

hashCode() 3 ns

Point creation 50 ns

So object creation is: ~ 47 ns

Thread creation: 700 ns



```
Mark6("Thread create start",  
    i -> {  
        Thread t = new Thread(() -> {  
            for (int j=0; j<1000; j++) //most iterations not done  
                ai.getAndIncrement(); // Why?  
        });  
        t.start();  
        return t.hashCode();  
    });
```

Takes ~ 47000 ns

- So, a lot of work goes into starting a thread
- Even after creating it
- Note: does not include executing the loop (why?)

Never create threads for small computations

Performance measurements: motivation and introduction

Pitfalls (and voiding them)

Calculating means and variance (efficiently)

Measurements of thread and lock overhead

Algorithms for parallel computing



Quicksort: <https://www.chrislaux.com/quicksort.html>

```
private static void qsort(int[] arr, int a, int b) {
    if (a < b) {
        int i = a, j = b;
        int x = arr[(i+j) / 2];
        do {
            while (arr[i] < x) i++;
            while (arr[j] > x) j--;
            if (i <= j) { swap(arr, i, j); i++; j--; }
        } while (i <= j);
        qsort(arr, a, j); qsort(arr, i, b);
    }
}
```

see SearchAndSort.java in week 03 material

Prime counting: <https://www.dcode.fr/prime-number-pi-count>


```
long count = 0;
final int from = 0, to = range;
for (int i=from; i<to; i++) if (isPrime(i)) count++;
```

see TestCountPrimes.java in week 03 material

Multi-threaded version of Quicksort



37



```
class Problem {  
    public int[] arr;  
    public int low, high;  
}
```

```
class ProblemHeap {  
    list<Problem> heap= new List<Problem>;  
    ... }
```

```
private static void qsort(Problem problem, ProblemHeap heap) {  
    int[] arr= problem.arr;  
    int a= problem.a;  
    int.b= problem.b;  
    ...  
    heap.add(new Problem(arr, a, j)); //qsort(arr, a, j);  
    heap.add(new Problem(arr, i, b)); //qsort(arr, i, b);  
}
```



quicksort 1	14196896.3	ns	136477.51
quicksort 2	8112412.2	ns	67791.32
quicksort 4	4912498.3	ns	71961.04
quicksort 8	3880639.1	ns	32812.31
quicksort 16	4553503.8	ns	40945.07
quicksort 32	6312270.0	ns	43905.97

Disappointing ?

Multithreaded version of CountPrimes

39



2, 3, 4, 5,



thread0



thread1

range



threadN

Code for exercises week03: `testCountPrimes.java`

Mark7 Count Primes



countSequential		5922958.0	ns	289879.33
countParallel	1	7107236.6	ns	448417.55
countParallel	2	6069944.7	ns	802224.61
countParallel	3	3621185.5	ns	152693.03
countParallel	4	3124067.0	ns	640480.51
countParallel	5	3699514.7	ns	364428.77
countParallel	6	4114074.2	ns	642562.19
countParallel	7	2049595.7	ns	26888.15
countParallel	8	1801465.6	ns	12532.85
countParallel	9	1793099.1	ns	11017.57
countParallel	10	1798921.4	ns	11541.43
countParallel	11	1807408.3	ns	9763.61

To be continued in week 6

