

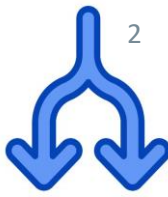


Practical Concurrent and Parallel Programming V

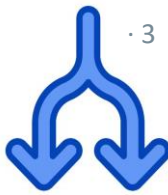
Parallel Streams

Raúl Pardo

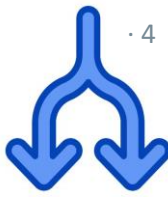
Exercise rooms



- Since rooms 2A12-14 have been sufficient for exercises in the past weeks, we have cancel the booking of the other two rooms.




- Oral feedback sessions are supposed to be fixed
 - Once you pick a slot it cannot change. Exceptionally some weeks you may negotiate alternative times with the TA/teacher you picked. But you should not change your choice in the scheduler.
 - Please note that for TAs/teachers to effectively prepare for exercise sessions we need stability and minimize changes
- Requests for written feedback have increased (too much)
 - Written feedback should be your last resource, you should use it only if it is not possible to attend oral feedback sessions.
 - We are open discuss alternative slots for oral feedback if this is the problem
 - Remember also that the exam will be like an oral feedback session. There is no better preparation for the exam than attending oral feedback sessions.



- I haven't forgotten about answering the comment on answering questions on the slides, it is coming this week
 - For this week, I already have a hidden version of the slides with answers to the questions

Previously on PCPP...




- Thread-safe classes
- Safe publication
- Immutability
- Instance confinement
- Synchronization primitives (synchronizers)
 - Semaphores
 - Barriers
- Producer-consumer problem

Agenda



- Data independence
- Lambda expressions
- Java Streams
- Parallel Java Streams



“Writing thread-safe code is, at its core, about managing access to shared mutable data”

Goetz

Data independence among threads



- Partitioning data so that threads only access disjoint memory is another way to ensure thread-safety

Data independence among threads



- Partitioning data so that threads only access disjoint memory is another way to ensure thread-safety

```
int parties          = 10;
CyclicBarrier cb     = new CyclicBarrier(parties);
int[] shared_array = new int[parties];
...
for (int i = 0; i < parties; i++) {
    new SetterClass(i).start();
}
...
public class SetterClass extends Thread {
    int index;
    public SetterClass(int index) {this.index = index;}

    public void run() {
        shared_array[index] = index+1;
        cb.await();
        // After this point the array is initialized and it is safe to read it
    }
}
```

Remember this example
from last week

Data independence among threads

· 10



- Partitioning data so that threads only access disjoint memory is another way to ensure thread-safety

```
int parties          = 10;
int[] shared_array = new int[parties];
...
for (int i = 0; i < parties; i++) {
    new SetterClass(i).start();
}
...
public class SetterClass extends Thread {
    int index;
    public SetterClass(int index) {this.index = index;}

    public void run() {
        // If the thread only works on share_array[index] the program is thread-safe
        shared_array[index] = index+1;
    }
}
```

If threads only access disjoint regions of the array the program is thread-safe

Data independence among threads

· 10



- Partitioning data so that threads only access disjoint memory is another way to ensure thread-safety

```
int parties          = 10;
int[] shared_array = new int[parties];
...
for (int i = 0; i < parties; i++) {
    new SetterClass(i).start();
}
...
public class SetterClass extends Thread {
    int index;
    public SetterClass(int index) {this.index = index;}

    public void run() {
        // If the thread only works on share_array[index] the program is thread-safe
        shared_array[index] = index+1;
    }
}
```

If threads only access disjoint regions of the array the program is thread-safe

BTW, this is the basis for GPU computing

Data independence among threads



· 11

- Partitioning data so that threads only access disjoint memory is another way to ensure thread-safety
- Java does not have intrinsic mechanisms to ensure this; it depends on the programmer

Data independence among threads



· 11

- Partitioning data so that threads only access disjoint memory is another way to ensure thread-safety
- Java does not have intrinsic mechanisms to ensure this; it depends on the programmer
- However, some parts of the standard library may help us
 - ThreadLocal
 - Parallel Streams
 - ...

Lambda Expressions



- One argument
 - `Function<Integer, Integer> f = (x) -> x+1`
 - $f : \mathbb{N} \rightarrow \mathbb{N} \mid f(x) = x + 1$
 - $f(1) \leftrightarrow \text{f.apply}(1)$
- Two arguments
 - `BiFunction<Integer, Integer, Integer> f = (x,y) -> x+y`
 - $f : \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N} \mid f(x, y) = x + y$
 - $f(1,2) \leftrightarrow \text{f.apply}(1,2)$
- ... (you see where we are going)



- Zero arguments
 - `Supplier<Integer> f = () -> 2`
 - $f : \mathbb{N} \mid f() = 2$
 - $f() \leftrightarrow f.get()$
- No return type (`void`)
 - `Consumer<Integer> f = (x) -> System.out.println(x);`
 - $f : \mathbb{N} \rightarrow Unit \mid f(x) = ???$
 - $f(2) \leftrightarrow f.accept(2)$



- Zero arguments

- `Supplier<Integer> f = () -> 2`
- $f : \mathbb{N} \mid f() = 2$
- $f() \leftrightarrow f.get()$

What about zero arguments and no return type?
For instance,

```
() -> System.out.println("Hej")
```

- No return type (`void`)

- `Consumer<Integer> f = (x) -> System.out.println(x);`
- $f : \mathbb{N} \rightarrow Unit \mid f(x) = ???$
- $f(2) \leftrightarrow f.accept(2)$



- The methods of an object may also be referenced as follows
- **Class::method**
 - `BiFunction<String,Integer,Character> f = String::charAt`
 - `f.apply(s,i) <-> s.charAt(i)`
 - `Function<Person,String> f = Person::getName`
 - `System.out::println`
 - ...
- **<Object instance>::method**
 - `Function<Integer,Character> f = "01234"::charAt`
 - `f.apply(i) <-> "01234".charAt(i)`

Higher-order functions



- A *higher-order* function is a function that
 - Takes functions as parameters, and/or
 - Returns a function
- Java streams support higher-order functions
 - map, reduce, filter, etc...

Java Streams



- A Java stream is a finite or infinite sequence of Objects
- Java streams use lazy evaluation
- The execution of operations over Java streams is typically efficient
- Operations on java streams may be easily executed parallel



- A Java stream is a finite or infinite sequence of Objects
- Java streams use lazy evaluation
- The execution of operations over Java streams is typically efficient
- Operations on java streams may be easily executed parallel

Why can operations be easily parallelized?

Defining a Java stream



- Using the `Arrays` class
 - `Arrays.stream(array)`
- Most Java collections have a method `stream()` that turns the collection into a stream
- `Stream.of(1,2,3,4)` creates a stream with those elements
- Functional iterators for infinite streams
 - `IntStream nats = IntStream.iterate(0, x->x+1)`
- `Stream::generate`
 - See file `StreamExample.java`
- `BufferedReader` (important for exercises)

```
Stream<String>
```

```
lines()
```

Returns a `Stream`, the elements of which are lines read from this `BufferedReader`

Intermediate vs Terminal operations



- Operations on Java streams may be
 - Intermediate
 - Terminal
- Intermediate operations return a new stream
 - Symbolic; they are recorded but the computation is postponed
- Terminal operations return a value
 - Execute the pipeline of intermediate operations to compute the value

- Examples of intermediate operations are:
 - filter – takes a lambda expression lambda returning a boolean, if the boolean is true the element is included in the output stream
 - map – takes a lambda computing a value. The output stream consists of computed values.
 - limit(n) – takes an integer and returns a stream of the first n elements
 - skip(n) – takes an integer and returns a stream without the first n elements
 - distinct – returns a stream without duplicated elements
 - sorted - returns a stream with the elements sorted



Intermediate operations

See Sestoft's Java precisely and
the java documentation for a
complete list





· 23

- Examples of intermediate operations are:
 - filter – takes a lambda expression lambda returning a boolean, if the boolean is true the element is included in the output stream
 - map – takes a lambda computing a value. The output stream consists of computed values.
 - limit(n) – takes an integer and returns a stream of the first n elements
 - skip(n) – takes an integer and returns a stream without the first n elements
 - distinct – returns a stream without duplicated elements
 - sorted - returns a stream with the elements sorted




map ($x \rightarrow x+2$)







Here it is important to note that the computation has not been executed only recorded




map ($x \rightarrow x+2$)





Intermediate operations | map & limit



Here it is important to note that the computation has not been executed only recorded




`map (x -> x+2)`




`limit (2)`





Intermediate operations | map & limit



Here it is important to note that the computation has not been executed only recorded



`map (x -> x+2)`



`limit (2)`



Is this function executed at this point?

- Remember, terminal operations take a stream and produce value
 - Not lazy; they execute the pipeline of intermediate operations to compute the value
- Java streams divides further terminal operations into:
 - Reduce
 - Collect



- Reduce all elements of the stream to a single value by applying a function
- **reduce(identity, accumulator)**
 - **identity**: The identity element is both the initial value of the reduction and the default result if there are no elements in the stream.
 - **accumulator**: The accumulator function takes two parameters: a partial result of the reduction and the next element of the stream.
- Example
 - Sum of squares of first 100 natural numbers
 - `IntStream.range(0,100).reduce(0, (a,b) -> a+b*b)`
- Implementation

```
T result = identity;
for (T element : this Stream)
    result = accumulator.apply(result, element)
return result;
```




- Reduce can also be called without identity parameter
- Then it returns an **Optional** value
 - A container object which may or may not contain a non-null value.
 - Needed in case the reduction is performed on an empty stream.
- Example
 - Sum of squares of first 100 natural numbers
 - `IntStream.range(0,100).reduce((a,b) -> a+b*b).orElse(0)`
- There exist other built-in reductions: sum, max, min, average, etc...

Example with everything so far



- Here is an example with everything we have seen so far
 - Amount of even numbers in the range 0 to 99

```
IntStream.iterate(0,x->x+1)
    .limit(100)
    .filter(x -> x%2==0)
    .map(x -> 1)
    .reduce(0, (a,b) -> a+b) ;
```

Example with everything so far



· 30

- Here is an example with everything we have seen so far
 - Amount of even numbers in the range 0 to 99

```
IntStream.iterate(0,x->x+1)
    .limit(100)
    .filter(x -> x%2==0)
    .map(x -> 1)
    .reduce(0, (a,b) -> a+b) ;
```


A sequence of stream
method calls is commonly
known as *stream pipeline*

Example with everything so far



0	1	2	3	...
---	---	---	---	-----

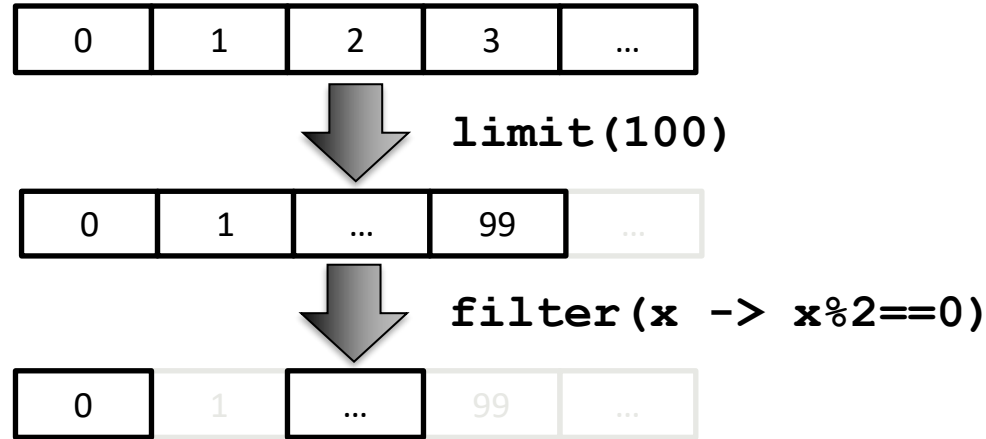
Example with everything so far



Example with everything so far



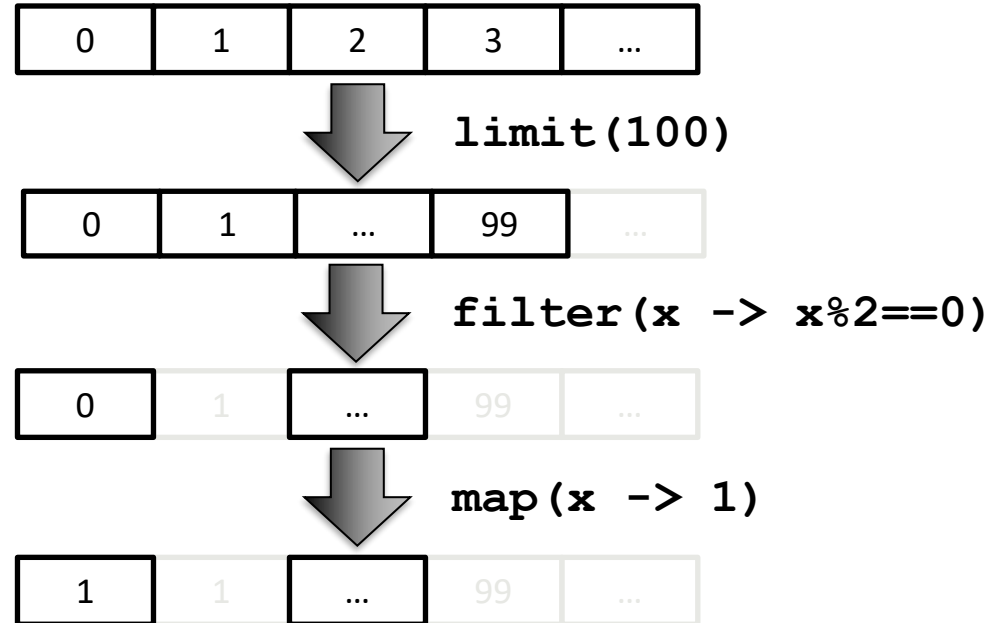
· 31



Example with everything so far



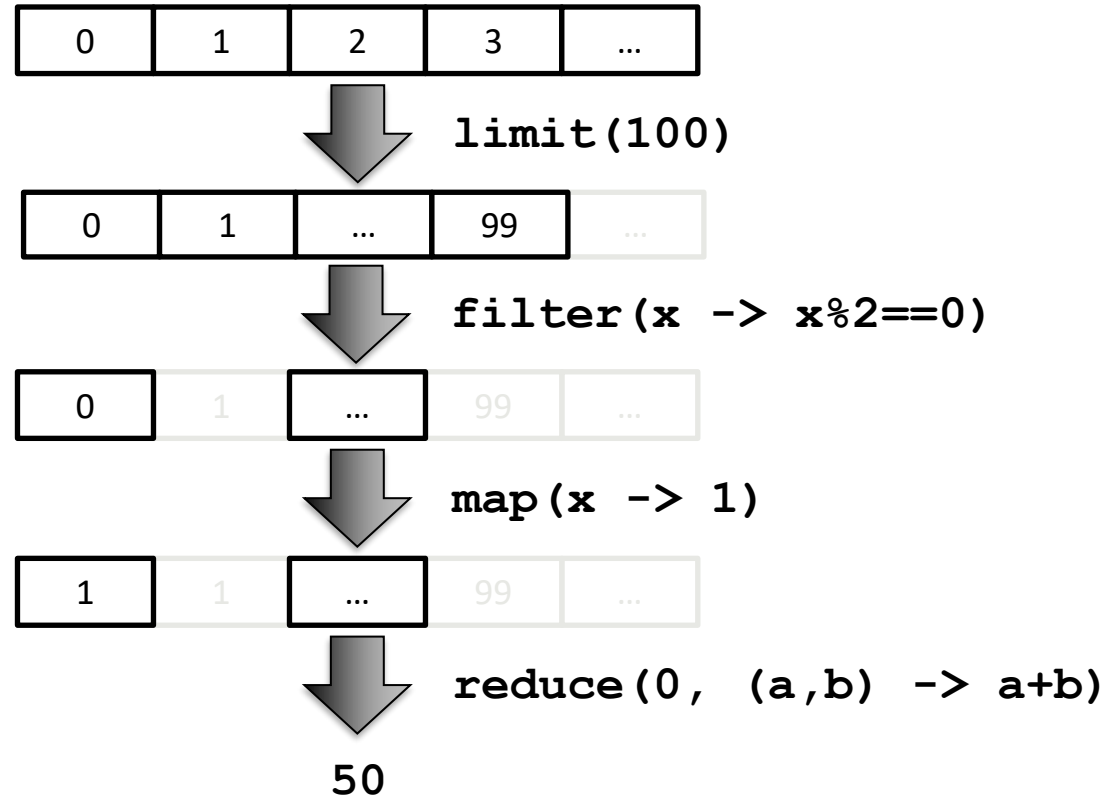
· 31



Example with everything so far



· 31





- It is a reduction operation that allows to collect the results of a stream into a Java collection or summarize them using complex criteria
- For instance, converting a stream into a list

```
List<Integer> l = randomEmployees()  
    .limit(50)  
    .map(Employee::getId)  
    .collect(Collectors.toList());
```

See `StreamExample.java`



- It is a reduction operation that allows to collect the results of a stream into a Java collection or summarize them using complex criteria
- For instance, converting a stream into a list

```
List<Integer> l = randomEmployees()  
    .limit(50)  
    .map(Employee::getId)  
    .collect(Collectors.toList());
```

Is this list symbolic like a stream or are values evaluated and stored in memory?

See `StreamExample.java`

- groupingBy is a special type of collector returning a Map where
 - Keys are generated based on some grouping criteria
 - Values are lists of elements (or operations on these lists of elements) of the stream matching the key

- Group employees by department


```
Map<String,List<Employee>> m = randomEmployees()  
    .limit(50)  
    .collect(Collectors.groupingBy(Employee::getDept));
```

Id: 0 Dept: CS Salary: 151	Id: 1 Dept: BI Salary: 150	Id: 3 Dept: DD Salary: 149	Id: 0 Dept: DD Salary: 10	...
----------------------------------	----------------------------------	----------------------------------	---------------------------------	-----

- Group employees by department

```
Map<String,List<Employee>> m = randomEmployees()  
    .limit(50)  
    .collect(Collectors.groupingBy(Employee::getDept));
```

Id: 0 Dept: CS Salary: 151	Id: 1 Dept: BI Salary: 150	Id: 3 Dept: DD Salary: 149	Id: 0 Dept: DD Salary: 10	...
----------------------------------	----------------------------------	----------------------------------	---------------------------------	-----



`collect(Collectors.groupingBy(Employee::getDept))`


CS	Id: 0 Dept: CS Salary: 151	...	
BI	Id: 1 Dept: BI Salary: 150	...	
DD	Id: 3 Dept: DD Salary: 149	Id: 0 Dept: DD Salary: 10	...

Is this map symbolic like a stream or are values evaluated and stored in memory?

- Group employees by department

```
Map<String,List<Employee>> m = randomEmployees()  
    .limit(50)  
    .collect(Collectors.groupingBy(Employee::getDept));
```

Id: 0 Dept: CS Salary: 151	Id: 1 Dept: BI Salary: 150	Id: 3 Dept: DD Salary: 149	Id: 0 Dept: DD Salary: 10	...
----------------------------------	----------------------------------	----------------------------------	---------------------------------	-----



`collect(Collectors.groupingBy(Employee::getDept))`


CS	Id: 0 Dept: CS Salary: 151	...	
BI	Id: 1 Dept: BI Salary: 150	...	
DD	Id: 3 Dept: DD Salary: 149	Id: 0 Dept: DD Salary: 10	...



- Get total salary per employee

```
randomEmployees().limit(50)  
  .collect(Collectors.groupingBy(Employee::getId, Collectors.summingInt(Employee::getSalary)))
```

Id: 0 Dept: CS Salary: 151	Id: 1 Dept: BI Salary: 150	Id: 3 Dept: DD Salary: 149	Id: 0 Dept: DD Salary: 10	...
----------------------------------	----------------------------------	----------------------------------	---------------------------------	-----



```
collect(Collectors.groupingBy(Employee::getId),  
  Collectors.summingInt(Employee::getSalary))
```

0	Id: 0 Dept: CS Salary: 151	Id: 0 Dept: DD Salary: 10	...
1	Id: 1 Dept: BI Salary: 150	...	
3	Id: 3 Dept: DD Salary: 149	...	
...			



0	161
1	150
3	149
...	

Parallelization of Streams



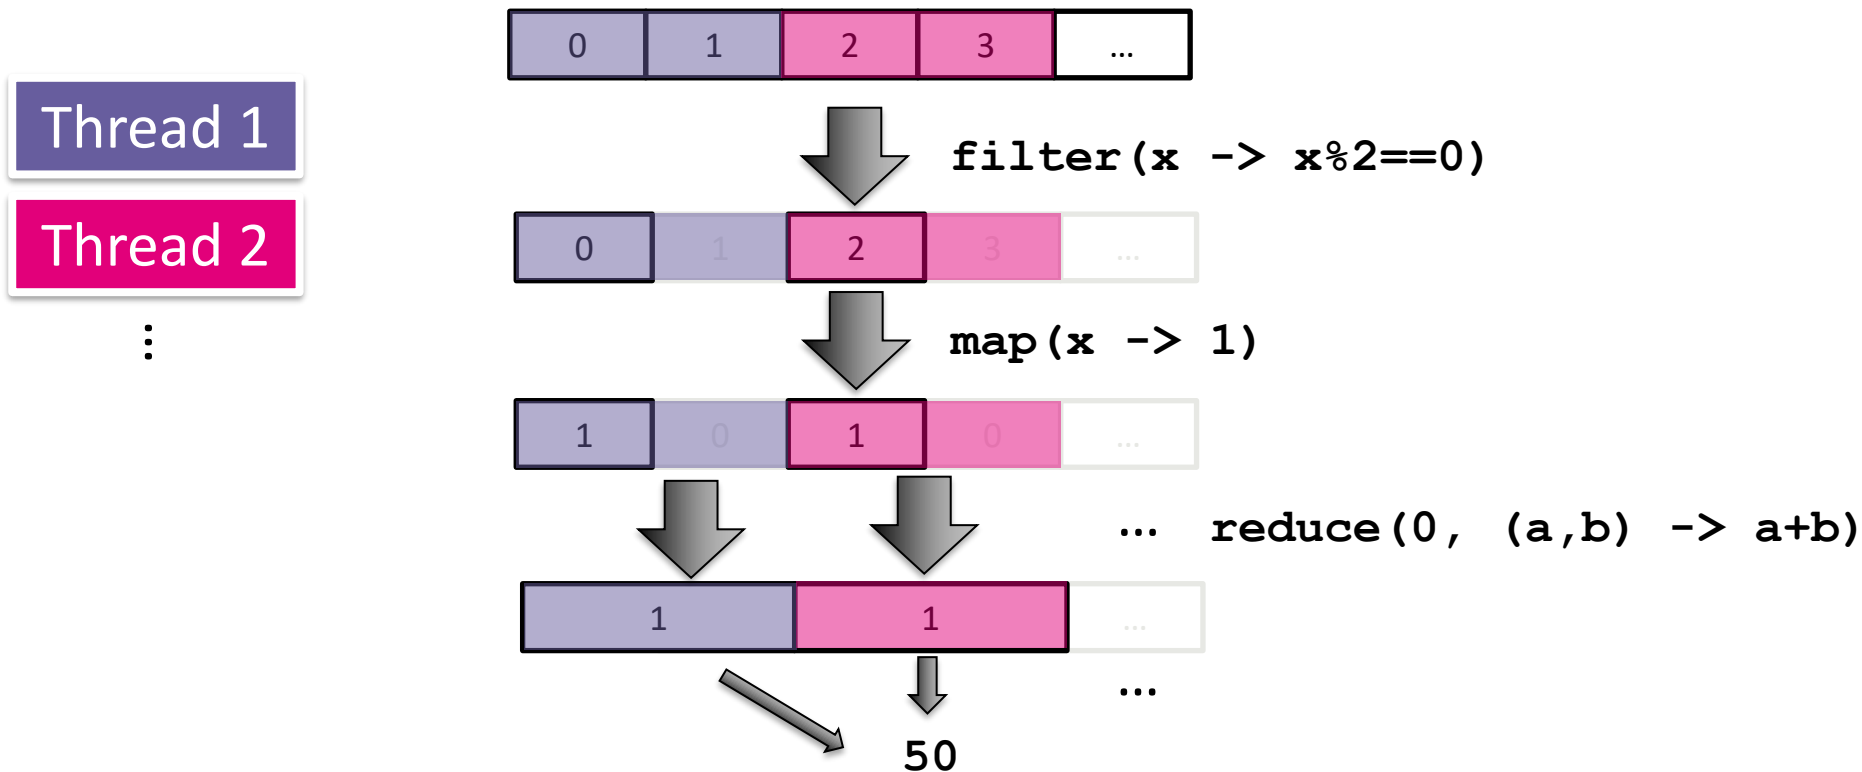
- You can create a parallel stream by calling
 - `parallelStream()` on, e.g., a collection, or
 - `parallel()` on a stream

Java Parallel Streams

· 40



- Parallelization of streams is very easy (remember the beginning of the lecture). Disjoint streams (from the original stream) are assigned to distinct threads from a thread pool (next lecture)



- Since execution is parallel the processing of the stream is not guaranteed to in order
- For instance, run this program
 - `IntStream.range(0,10).parallel().forEach(System.out::println);`
- In this case, it may be mitigated with `forEachOrdered`
 - `IntStream.range(0,10).parallel().forEachOrdered(System.out::println);`

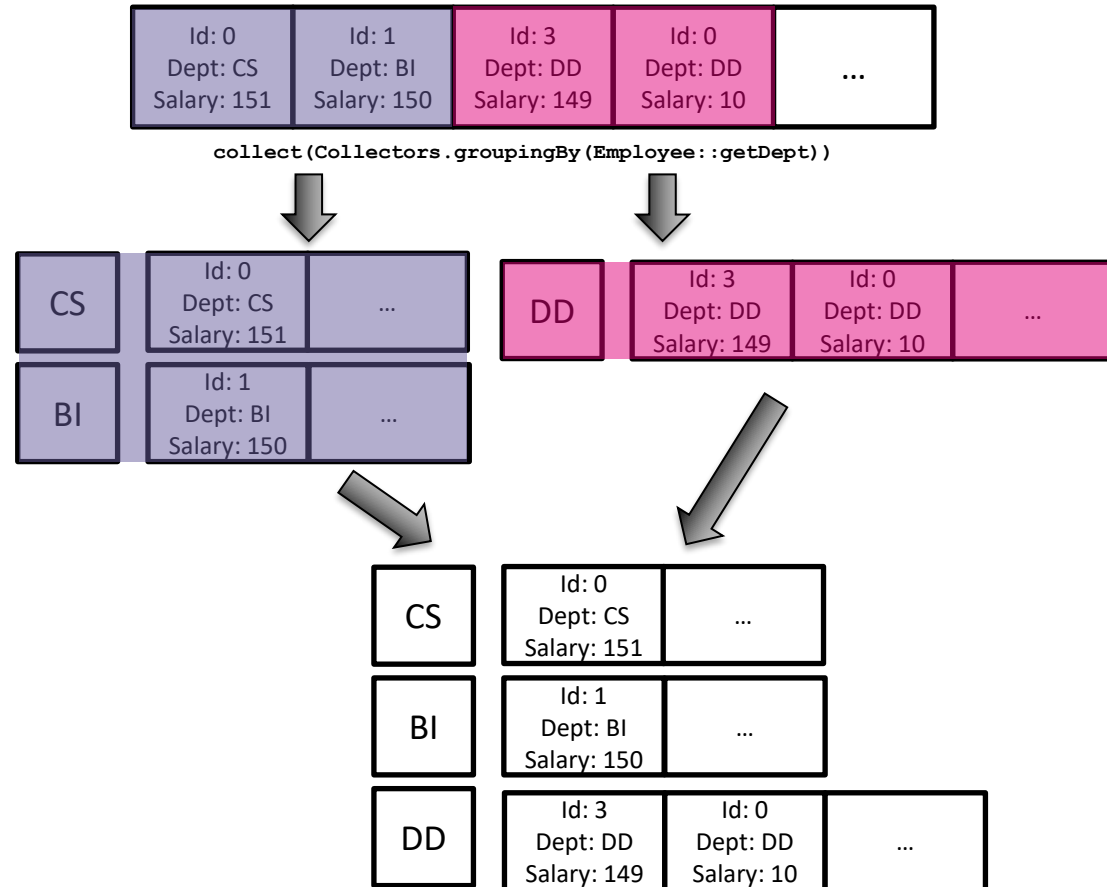
Would it be ok if parallel streams work concurrently on the Map resulting from a `groupingBy` operation?

Java Parallel Streams | Grouping

· 44



- In the local reduce operation, each thread creates locally its own “container” (collection or value) for the partial results
- In a final step, a single thread merges the local containers, so no race conditions arise
- However this step is very memory waste full, as it needs to create one container per partition of stream



- A possible solution is to use `groupingByConcurrent`
- This function uses a `ConcurrentHashMap` that all threads can access concurrently.

```
collect(Collectors.groupingByConcurrent(Employee::getDept))
```



- If you try to modify a stream you are operating you will get a `ConcurrentModificationException` at runtime
 - So don't do it 😊
- Cannot be detected a compile time. It depends on the programmer
- From the Java documentation
 - *Streams enable you to execute possibly-parallel aggregate operations* over a variety of data sources, including even non-thread-safe collections such as `ArrayList`. *This is possible only if we can prevent interference with the data source during the execution of a stream pipeline.* [...] For most data sources, preventing interference means ensuring that the data source is not modified at all during the execution of the stream pipeline.

Counting primes on Java 8 streams

47



Our old standard Java for loop:

Classical efficient
imperative loop

```
int count = 0;
for (int i=0; i<range; i++)
    if (isPrime(i))
        count++;
```

Sequential Java 8 stream:

```
IntStream.range(0, range)
    .filter(i -> isPrime(i))
    .count()
```

Parallel Java 8 stream:

```
IntStream.range(0, range)
    .parallel()
    .filter(i -> isPrime(i))
    .count()
```

PCPP 2019

Counting primes on Java 8 streams

47



Our old standard Java for loop:

Classical efficient
imperative loop

```
int count = 0;
for (int i=0; i<range; i++)
    if (isPrime(i))
        count++;
```

Sequential Java 8 stream:

Pure functional
programming ...

```
IntStream.range(0, range)
    .filter(i -> isPrime(i))
    .count()
```

Parallel Java 8 stream:

```
IntStream.range(0, range)
    .parallel()
    .filter(i -> isPrime(i))
    .count()
```

PCPP 2019

Counting primes on Java 8 streams

47



Our old standard Java for loop:

Classical efficient
imperative loop

```
int count = 0;
for (int i=0; i<range; i++)
    if (isPrime(i))
        count++;
```

Sequential Java 8 stream:

Pure functional
programming ...

```
IntStream.range(0, range)
    .filter(i -> isPrime(i))
    .count()
```

Parallel Java 8 stream:

... and thus
parallelizable and
thread-safe

```
IntStream.range(0, range)
    .parallel()
    .filter(i -> isPrime(i))
    .count()
```

PCPP 2019

Performance results (!!)



Counting the primes in 0 ...99,999

Method	Intel i7 (ms)	AMD Opteron (ms)
Sequential for-loop	9.9	40.5
Sequential stream	9.9	40.8
Parallel stream	2.8	1.7

Functional streams give the simplest solution

Nearly as fast as tasks and threads, or faster:

Intel i7 (4 cores) speed-up: 3.6 x

AMD Opteron (32 cores) speed-up: 24.2 x

The future is parallel – and functional 😊

Agenda



- Data independence
- Lambda expressions
- Java Streams
- Parallel Java Streams